

Transactions

事务：保证多个操作的原子性

特性：所有命令都在同一台服务器上执行

To do transactions in Jedis, you have to wrap operations in a transaction block, very similar to pipelining: 事务块

```
jedis.watch (key1, key2, ...); 先加锁
Transaction t = jedis.multi();
t.set("foo", "bar");
t.exec();
```

Note: when you have any method that returns values, you have to do like this:

```
Transaction t = jedis.multi();
t.set("foo", "bar");
Response<String> result1 = t.get("foo");

t.zadd("foo", 1, "barowitch"); t.zadd("foo", 0, "barinsky"); t.zadd("foo", 0,
Response<Set<String>> sose = t.zrange("foo", 0, -1); // get the entire sorte
t.exec(); // dont forget it

String foolbar = result1.get(); // use Response.get() to
int soseSize = sose.get().size(); // on sose.get() you ca

// List<Object> allResults = t.exec(); // you could still get all
```

Note that a Response Object does not contain the result before t.exec() is called (it is a kind of a Future). Forgetting exec gives you exceptions. In the last lines, you see how transactions/pipelines were dealt with before version 2. You can still do it that way, but then you need to extract objects from a list, which contains also Redis status messages.

Note 2: Redis does not allow to use intermediate results of a transaction within that same transaction. This does not work:

```
    这样不起作用！
// this does not work! Intra-transaction dependencies are not supported by Red
jedis.watch(...);
Transaction t = jedis.multi();
if(t.get("key1").equals("something"))
    t.set("key2", "value2");
else
```

```
t.set("key", "value");
```

However, there are some commands like `setnx`, that include such a conditional execution. Those are of course supported within transactions. You can build your own customized commands using `eval/ LUA scripting`.

构建自定义的命令

Pipelining 管道：执行一连串不同的命令

Sometimes you need to send a bunch of different commands. A very cool way to do that, and have better performance than doing it the naive way, is to use `pipelining`. This way you send commands without waiting for response, and you actually read the responses at the end, which is faster.

使用管道有更好的性能

以这种方式发送命令，无需等待响应，同时会在最后才真正读取响应信息，这是非常快的。

Here is how to do it:

特性：各个命令可以被发送到不同的服务器，不保证所有命令的原子性

```
Pipeline p = jedis.pipelined();
p.set("foo", "bar");
p.zadd("foo", 1, "barowitch"); p.zadd("foo", 0, "barinsky"); p.zadd("foo", 0,
Response<String> pipeString = p.get("foo");
Response<Set<String>> sose = p.zrange("foo", 0, -1);
p.sync();

int soseSize = sose.get().size();
Set<String> setBack = sose.get();
```

For more explanations see code comments in the transaction section.

Publish/Subscribe 消息发布/订阅

To subscribe to a channel in Redis, create an instance of `JedisPubSub` and call subscribe on the Jedis instance:

要想在Redis中订阅一个通道，需要创建一个`JedisPubSub`实例，并调用`subscribe`方法。

```
class MyListener extends JedisPubSub {
    public void onMessage(String channel, String message) {
    }

    public void onSubscribe(String channel, int subscribedChannels) {
    }

    public void onUnsubscribe(String channel, int subscribedChannels) {
    }

    public void onPSubscribe(String pattern, int subscribedChannels) {
    }
}
```

```

    }

    public void onPUnsubscribe(String pattern, int subscribedChannels) {
    }

    public void onPMessage(String pattern, String channel,
        String message) {
    }
}

MyListener l = new MyListener();

jedis.subscribe(l, "foo");

```

注意：订阅是一个阻塞的操作，因为它会轮询Redis线程的响应信息！

Note that subscribe is a blocking operation because it will poll Redis for responses on the thread that calls subscribe. A single JedisPubSub instance can be used to subscribe to multiple channels. You can call subscribe or psubscribe on an existing JedisPubSub instance to change your subscriptions.

ShardedJedis 数据分片的Redis实现

Motivation 设计初衷

在标准的"主-从"方式中，一个主库负责写入请求，多个从库负责读请求。

In the normal Redis master-slave approach, generally there is one master that serves write requests, and many slaves that serve read requests. This means, the user has to take care of effectively distributing the load on the slaves. Furthermore, only reads scale with the number of slaves, but writes do not, since there can be only one master! With ShardedJedis you achieve scalability for both reads and writes. Sharding uses a technique called "consistent hashing" and assigns the keys equally on a set of redis servers according to some hash algorithm (md5 and murmur, the latter being less standard, but faster). A node like this is then called a "shard". A further advantage is that each shards only needs to have RAM 1/n the size of the total dataset (for n being the number of participating slaves).

使用ShardedJedis，可以实现可扩展的读取和写入操作。
分片(Sharding)使用了一种被称为"一致性哈希算法"的技术
更大的优势：每个分片仅需要总数据集大小的1/n内存(RAM)

The downside 缺点

Since each shard is a separate master, sharding has limited functionality: i.e. you cannot use transactions, pipelining, pub/sub, especially not across shards! However, generally it is feasible to do a not allowed operation, as long as the concerned keys are on the same shard (check / ask the forum). You can influence which key go to which shard by keytags (see below). A further downside is that in the current standard implementation, shards cannot be added or removed from a running

因为每个分片都是一个独立的主库，**分片是有功能限制的**：不能使用事务、管道、发布/订阅，**尤其是不能跨越多个分片！**

更大的**缺点**：分片不能从一个运行的ShardedJedis实例中动态地添加或删除

ShardedJedis. If you need this feature, there is an experimental reimplementation of ShardedJedis which allows adding and removing shards of a running ShardedJedis: [yaourt - dynamic sharding implementation](#) 看过源码了，还是需要手动去触发的。

General Usage: 一般用法：

1. Define your shards:

```
List<JedisShardInfo> shards = new ArrayList<JedisShardInfo>();
JedisShardInfo si = new JedisShardInfo("localhost", 6379);
si.setPassword("foobared");
shards.add(si);
si = new JedisShardInfo("localhost", 6380);
si.setPassword("foobared");
shards.add(si);
```

Then, there are two ways of using ShardedJedis. Direct connections or by using ShardedJedisPool. For reliable operation, the latter has to be used in a multithreaded environment. 出于**操作可靠性**考虑，在**多线程**环境中应该使用ShardedJedisPool

2.a) Direct connection:

```
ShardedJedis jedis = new ShardedJedis(shards);
jedis.set("a", "foo");
jedis.disconnect;
```

2.b) **Pooled connection**:

```
ShardedJedisPool pool = new ShardedJedisPool(new Config(), shards);
ShardedJedis jedis = pool.getResource();
jedis.set("a", "foo");
.... // do your work here
pool.returnResource(jedis); 需要显示地返回资源，这样设计不是很好！
.... // a few moments later
ShardedJedis jedis2 = pool.getResource();
jedis.set("z", "bar");
pool.returnResource(jedis);
pool.destroy();
```

3. **Disconnect / returnResource** returnResource应该在处理完业务逻辑后，立马就被调用

`pool.returnResource` should be called as soon as you are finished using jedis in a particular moment. If you don't, the pool may get slower after a while. `getResource`

getResource和returnResource操作是很快，因为没有新的连接需要被创建！但连接池的创建和释放是很慢的，因为这些都是真实的网络连接。

and `returnResource` are fast, since no new connection have to be created. Creation and destruction of a pool are slower, since these are the actual network connections. Forgetting pool.destroy keeps the connection open until timeout is reached. 当忘记调用`pool.destroy()`方法时，连接会一直开着，直到超时时间达到才会被关闭。

Determine information of the shard of a particular key

```
ShardInfo si = jedis.getShardInfo(key);
si.getHost/getPort/getPassword/getTimeout/getName
```

Force certain keys to go to the same shard 强制特定的键s落到同一个分片里

What you need is something called "keytags", and they are supported by Jedis. To work with keytags you just need to set a pattern when you instance ShardedJedis. For example:

```
ShardedJedis jedis = new ShardedJedis(shards,
ShardedJedis.DEFAULT_KEY_TAG_PATTERN);
```

You can create your own pattern if you want. The default pattern is {}, this means that whatever goes inside curly brackets will be used to determine the shard.

So for example:

```
jedis.set("foo{bar}", "12345");
```

and

```
jedis.set("car{bar}", "877878");
```

will go to the same shard.

Mixed approach "一主多从"架构：无法扩展写(Master)操作

If you want easy load distribution of ShardedJedis, but still need transactions/pipelining/pubsub etc, you can also mix the normal and the sharded approach: define a master as normal Jedis, the others as sharded Jedis. Then make all the shards slaveof master. In your application, direct your write requests to the master, the read requests to ShardedJedis. Your writes don't scale anymore, but you gain good read distribution, and you have transactions/pipelining/pubsub simply 写入是无法扩展的，但可以获得好的读取分发

using the master. Dataset should fit in RAM of master. Remember that you can improve performance of the master a lot, if you let the slaves do the persistence for the master! 记住：如果你使所有从库都与主库保持一致，可以极大地改善主库的性能！

Redis Cluster Redis集群

Sometime later 2011, there will be first versions of "redis cluster" which will be a much improved Sharded Jedis and should give back some if not all of the Redis functionalities you cannot have with shardedJedis. If you want to know more about redis cluster, youtube has a presentation of Salvatore Sanfilippo (the creator of Redis).

Monitoring 监控

To use the monitor command you can do something like the following:

```
new Thread(new Runnable() {
    public void run() {
        Jedis j = new Jedis("localhost");
        for (int i = 0; i < 100; i++) {
            j.incr("foobared");
            try {
                Thread.sleep(200);
            } catch (InterruptedException e) {
            }
        }
        j.disconnect();
    }
}).start();

jedis.monitor(new JedisMonitor() {
    public void onCommand(String command) {
        System.out.println(command);
    }
});
```

Misc

关于字符串和二进制的说明 - 哪个是原生的？

A note about String and Binary - what is native?

Redis/Jedis talks a lot about Strings. And here <http://redis.io/topics/internals> it says Strings are the basic building block of Redis. However, this stress on strings may be misleading. Redis' "String" refer to the C char type (8 bit), which is incompatible with Redis字符串引用C的字符字节(8位)，它与Java的字符串(16位)并不兼容

Redis仅识别预先定义的8位长度的数据块，所以正常情况下它是没法被转义为数据的(即"二进制安全")

Java Strings (16-bit). Redis sees only 8-bit blocks of data of predefined length, so normally it doesn't interpret the data (it's "binary safe"). Therefore in Java, byte[] data is "native", whereas Strings have to be encoded before being sent, and decoded after being retrieved by the SafeEncoder. This has some minor performance impact. In short: if you have binary data, don't encode it into String, but use the binary versions.

换句话说：如果你拥有二进制数据，不用将它编码成字符串，而是直接使用二进制版本。

A note on Redis' master/slave distribution

A Redis network consists of redis servers, which can be either masters or slaves. Slaves are synchronized to the master (master/slave replication). However, master and slaves look identical to a client, and slaves do accept write requests, but they will not be propagated "up-hill" and could eventually be overwritten by the master. It makes sense to route reads to slaves, and write demands to the master.

Furthermore, being a slave doesn't prevent from being considered master by another slave. 从库之间是对等的

读写分离：路由读请求到从库s，写入请求命令到主库，这是有意义的。